

# Principles of Programming

## CM10227

### Lecture S.1.: Introduction to Linux/Unix



Dr. Marina De Vos  
University of Bath  
Ext: 5053

Academic Year 2012-2013



# Outline

## 1 Shell Basics



# Resources

- Unix for Beginners. Dirk Vermeir
- <http://osl.iu.edu/~lums/swc/www/index.html>



# Outline

- 1 Shell Basics
  - The Shell
  - The File System
  - Some Linux Commands



# Why Command-line?

- Most modern tools have a graphical user interface (GUI)
  - Because they're easier to use
- But command-line user interfaces (CLUIs) still have their place
  - Easier (faster) to build new CLUI tools
    - Building a GUI takes time
    - Building a good GUI takes a lot of time
  - Higher action-to-keystroke ratio
    - Once you're over the (steeper) learning curve
  - Easier to see and understand what the computer is doing on your behalf
    - Which is part of what this course is about
  - Most important: it's easier to combine CLUI tools than GUI tools
    - Small tools, combined in many ways, can be very powerful



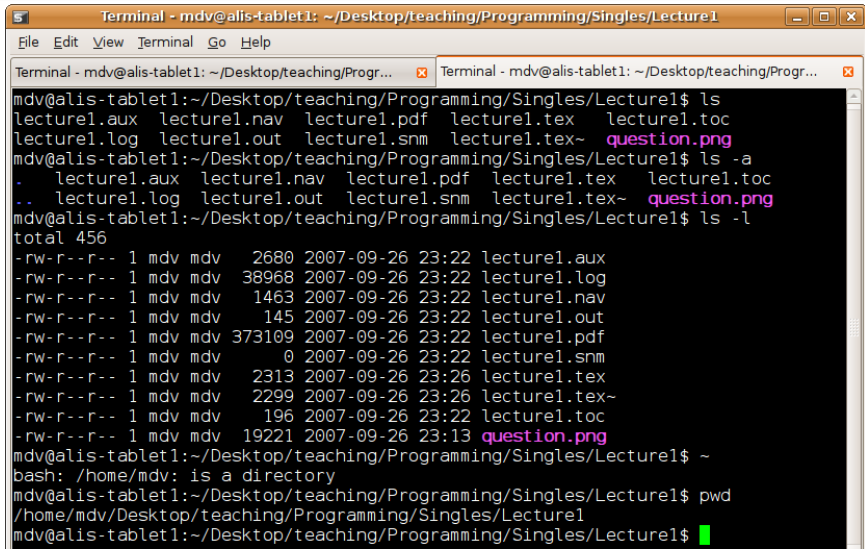
# The Shell

The most important command-line tool is the command shell (often just called the shell)

- Manages a user's interactions with the operating system by:
  - Reading commands from the keyboard
  - Figuring out what programs the user wants to run
  - Running those programs
  - Displaying their output on the screen
- Looks (and works) like an interactive terminal circa 1980



# The Terminal



A screenshot of a Linux terminal window titled "Terminal - mdv@alis-tablet1: ~/Desktop/teaching/Programming/Singles/Lecture1". The window has a menu bar with "File", "Edit", "View", "Terminal", "Go", and "Help". There are two tabs open, both showing the same path. The terminal content shows the user running several commands to list files in the current directory. The first command is `ls`, which lists files including `question.png` in pink. The second command is `ls -a`, showing hidden files like `lecture1.log`. The third command is `ls -l`, which displays a detailed file listing with permissions, sizes, dates, and file names. The output of `ls -l` shows files like `lecture1.aux`, `lecture1.log`, `lecture1.nav`, `lecture1.out`, `lecture1.pdf`, `lecture1.snm`, `lecture1.tex`, `lecture1.tex~`, `lecture1.toc`, and `question.png`. The user then enters `~`, which results in the message `bash: /home/mdv: is a directory`. Finally, the user enters `pwd`, which outputs `/home/mdv/Desktop/teaching/Programming/Singles/Lecture1`. The prompt `mdv@alis-tablet1:~/Desktop/teaching/Programming/Singles/Lecture1$` is shown at the bottom with a green cursor.

```
mdv@alis-tablet1:~/Desktop/teaching/Programming/Singles/Lecture1$ ls
lecture1.aux lecture1.nav lecture1.pdf lecture1.tex lecture1.toc
lecture1.log lecture1.out lecture1.snm lecture1.tex~ question.png
mdv@alis-tablet1:~/Desktop/teaching/Programming/Singles/Lecture1$ ls -a
. lecture1.aux lecture1.nav lecture1.pdf lecture1.tex lecture1.toc
.. lecture1.log lecture1.out lecture1.snm lecture1.tex~ question.png
mdv@alis-tablet1:~/Desktop/teaching/Programming/Singles/Lecture1$ ls -l
total 456
-rw-r--r-- 1 mdv mdv 2680 2007-09-26 23:22 lecture1.aux
-rw-r--r-- 1 mdv mdv 38968 2007-09-26 23:22 lecture1.log
-rw-r--r-- 1 mdv mdv 1463 2007-09-26 23:22 lecture1.nav
-rw-r--r-- 1 mdv mdv 145 2007-09-26 23:22 lecture1.out
-rw-r--r-- 1 mdv mdv 373109 2007-09-26 23:22 lecture1.pdf
-rw-r--r-- 1 mdv mdv 0 2007-09-26 23:22 lecture1.snm
-rw-r--r-- 1 mdv mdv 2313 2007-09-26 23:26 lecture1.tex
-rw-r--r-- 1 mdv mdv 2299 2007-09-26 23:26 lecture1.tex~
-rw-r--r-- 1 mdv mdv 196 2007-09-26 23:22 lecture1.toc
-rw-r--r-- 1 mdv mdv 19221 2007-09-26 23:13 question.png
mdv@alis-tablet1:~/Desktop/teaching/Programming/Singles/Lecture1$ ~
bash: /home/mdv: is a directory
mdv@alis-tablet1:~/Desktop/teaching/Programming/Singles/Lecture1$ pwd
/home/mdv/Desktop/teaching/Programming/Singles/Lecture1
mdv@alis-tablet1:~/Desktop/teaching/Programming/Singles/Lecture1$
```

# The Shell vs. the Operation System

- The shell is just one program among many
  - Many different ones have been written
  - **sh** was the first for Unix
    - Most others extend its capabilities in various ways
    - Which means that it's the lowest common denominator you can always rely on
  - We will use **bash** (the Bourne again shell)
    - Available just about everywhere
    - Even on Windows (thanks to Cygwin)





# The Shell vs. the Operation System

- In contrast, the operating system is not just another program
  - Automatically loaded when the computer boots up
  - The only program that can talk directly to the computer's hardware
    - I.e., read characters from the keyboard, or send drawing commands to the screen
  - Manages files and directories on the disk
  - Keeps track of who you are, and what you're allowed to do
  - You can run many instances of the shell on a computer at once, but it can only run one operating system at a time



# The File System

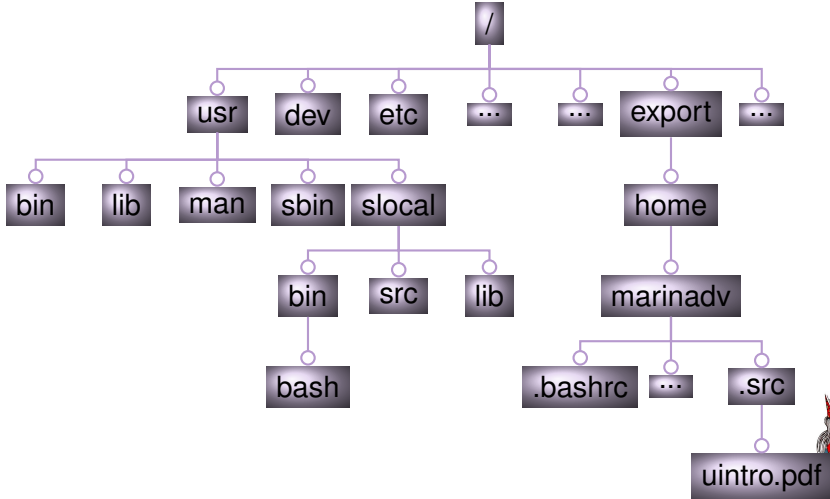
- The file system is the set of files and directories the computer can access

*Everything that stays put when you turn the computer off and restart it*

- Data is stored in files
  - By convention, files have two part names, like notes.txt or home.html
  - Most operating systems allow you to associate a filename extension with an application
    - E.g., .txt is associated with an editor, and .html with a web browser
  - But this is all just convention: you can call files (almost) anything you want
- Files are stored in directories (often called folders)
  - Directories can contain other directories, too
  - Results in the familiar directory tree



# Directory Tree



# Drives

- On Unix, the file system has a unique **root directory** called  
/
  - Every other directory is a child of it, or a child of a child, etc.
- On Windows, every **drive** has its own root directory
  - So C:\home\mdv\notes.txt is different from  
J:\home\mdv\notes.txt
  - When you're using Cygwin, you can also write C:\home\mdv  
as c:/home/mdv
  - Or as /cygdrive/c/home/mdv
    - Some Unix programs give ":" a special meaning, so Cygwin  
needed a way to write paths without it



# Paths

A **path** is a description of how to find something in a file system

- An **absolute path** describes a location from the root directory down
  - Equivalent to a street address
  - Always starts with "/"
  - E.g., /home/mdv is my home directory, and /courses/swc/lec/shell.swc is this file
- A **relative path** describes how to find something from some other location
  - Equivalent to saying, Four blocks north, and seven east
  - E.g., from /courses/swc, the relative path to this file is lec/shell.swc



# Special Paths

- Every program (including the shell) has a current working directory
  - Where am I?
  - Relative paths are deciphered relative to this location
  - It can change while a program is running
- Finally, two special names:
  - "." means the current directory
  - ".." means the directory immediately above this one
    - Also called the parent directory
    - In /courses/swc/data, .. is /courses/swc
    - In /courses/swc/data/elements, .. is /courses/swc/data



# File Systems

Most unix systems have several types of file systems

- Disk-based: UFS : to store all the files users create
- Network-based: NFS: to connect to (mount) drives outside the machine
- tmpfs file system: supports simulating a file system in main memory, possibly backed up by swap storage. This is ideal for temporary files for which fast access is important.
- swap: file system is used to provide backup storage for processes that must temporarily be swapped out
- proc file space: provides a file view on the attributes of processes



# pwd and ls

- **pwd** shows you the current directory

```
pwd
```

```
/home/mdv/Desktop/teaching/Programming/Singles/Lecture1
```

- **ls** shows you what's in the current directory

```
ls
```

```
lecture1.aux  lecture1.out  lecture1.tex  lecture1.vrb  
lecture1.log  lecture1.pdf  lecture1.tex~  question.png  
lecture1.nav  lecture1.sn  lecture1.toc  terminal.png
```





## More on ls

What actually happens when I type `ls` is:

- The operating system reads characters from the keyboard
- Passes them to the shell (because it's the currently active window on my desktop)
- The shell breaks the line of text it receives into words
- Looks for a program with the same name as the first word (i.e., the command to run)
  - Describe in a moment how the shell knows where to look
- Runs that program
- Reads the program's output and sends it back to the operating system for display



# Flags

- Flags are command-line option you can pass to commands
- Can tell ls to produce more informative output by giving it some flags
- By convention, flags start with "-", as in "-c" or "-l"
- For example: show directories with trailing slash

```
ls -F
```

```
bluej.png      code.sty~      copyright.tex~  rights.png      uintro.pdf
cm10192.tex    computer.jpg    Doubles/        Singles/
cm10192.tex~   computer.png    projects.zip     Stylefiles/
code.sty       copyright.tex*  python.png       template.tex~
```

- a: gives you all files starting with ".", which are normally hidden
- l: provides long listing format. provides permissions, size, latest access



# Finding your way

- man pages: provide an overview of the functionality of a command.
  - man ls
- apropos: provides all commands related to a certain topic
  - apropos(permissions)
- --help: provides support for a specific command
  - ls --help



# Manipulating Files and Directories

Lets work by example:

- let us create a temporary directory and play around in there

```
mkdir temp
```

- Note: no output
- The -v (verbose) flag tells mkdir to print a confirmation message
- Now go into that directory

```
cd temp
```

- Changes the shell's notion of our current working directory

```
pwd
```

```
/home/mdv/programming1/temp
```



## Manipulating Files and Directories II

- No files there yet:

```
ls -a
```

```
. ..
```

- Use the editor of your choice (emacs,vim) to create a file called earth.txt with the following contents:

```
Name:  Earth
Period: 365.26 days
Inclination: 0.00
Eccentricity: 0.02
Object: Planet
```



## Manipulating Files and Directories III

- Easiest way to create a similar file venus.txt is to copy the one we have

```
cp earth.txt venus.txt
```

```
ls -t
```

```
venus.txt    earth.txt
```

- Note: the -t option tells ls to list newest first
- Check the contents of the file using cat (short for concatenate)
- Just prints the contents of a file to the screen
- You can also use more or less



# Manipulating Files and Directories IV

- Edit the file so that looks like:
- Compare the sizes of the two files using `wc` (for word count) and Compare the two files using `diff`

```
Name: Venus
Period: 224.70 days
Inclination: 3.39
Eccentricity: 0.01
Object: Planet
```

```
wc earth.txt venus.txt
```

```
4  9 69 earth.txt
4  9 69 venus.txt
8 18 138 total
```

```
diff earth.txt venus.txt
```

```
1,4c1,4
< Name: Earth
< Period: 365.26 days
< Inclination: 0.00
< Eccentricity: 0.02
---
> Name: Venus
> Period: 224.70 days
> Inclination: 3.39
> Eccentricity: 0.01
```



# Manipulating Files and Directories V

- Linux does not care about filename extensions.
- `cp earth.txt earth.pdf` is valid although not a very sensible thing to do
- we can rename it using `mv earth.pdf earth2.txt`
- Removing a file can be done using `rm`, like for example `rm earth2.txt`
- A empty directory can be removed with `rmdir` or `rm -r` which recursively removed all files.





# Wildcards

- Some characters (**wildcards**) mean special things to the shell
  - \* matches zero or more characters
    - So `ls *.f77` lists all the Fortran-77 files in a directory

```
wc *.txt
 4   9   69 earth.txt
 4   9   69 venus.txt
 8  18  138 total
```

- ? matches any single character
  - So `ls ???.txt` lists all the text files with two-letter prefixes
  - And `ls ??*.*` lists all the files with two-letter prefixes, and any extension
  - `~` on its own means the users home directory
  - `harry` means Harry's home directory
- Note: the shell expands wildcards before running commands
- Note: Be careful using `rm` in conjunction with \*



# Users

- Users have a **user name** and a **password**
- a user also has a **home directory**, and a **shell program**.
- Internally, the system uses so-called **UID** numbers to identify users.
- All this information is stored in the file **/etc/passwd**
- This also stores the user primary group id (GID) identifying a group to which the user belongs.
- A **group** is an arbitrary set of users
- A user can belong to several groups
- `whoami`, `users`, `groups` provide you with information regarding users and groups
- There is one special user with UID 0, called **root**
- This user is often called the super user because he can access all resources on the system, independently of any specific permissions



# Ownership

- Each file has a user as **owner** and a group as **group owner**.
- Using `chmod` the owner can change permissions that determine the type of access (read, write or execute) ...
- allowed to three categories of users: the owner herself, the users belonging to the group owner group, and all other users.
- Note that "execute" permission on a directory is interpreted as "permission to traverse"

```
ls -l
-rw-r--r-- 1 mdv mdv      84 2007-09-27 23:08 earth.pdf
chmod g+w earth.txt
ls -l
-rw-rw-r-- 1 mdv mdv      84 2007-09-27 22:38 earth.txt
```

