

# Principles of Programming

## CM10227

Lecture S.1.: Introduction to Linux/Unix



Dr. Marina De Vos  
University of Bath  
Ext: 5053

Academic Year 2012-2013



# Outline

## 1 Shell Basics



# Resources

- Unix for Beginners. Dirk Vermeir
- <http://osl.iu.edu/~lums/swc/www/index.html>



# Why Command-line?

- Most modern tools have a graphical user interface (GUI)
  - Because they're easier to use
- But command-line user interfaces (CLUIs) still have their place
  - Easier (faster) to build new CLUI tools
    - Building a GUI takes time
    - Building a good GUI takes a lot of time
  - Higher action-to-keystroke ratio
    - Once you're over the (steeper) learning curve
  - Easier to see and understand what the computer is doing on your behalf
    - Which is part of what this course is about
  - Most important: it's easier to combine CLUI tools than GUI tools
    - Small tools, combined in many ways, can be very powerful



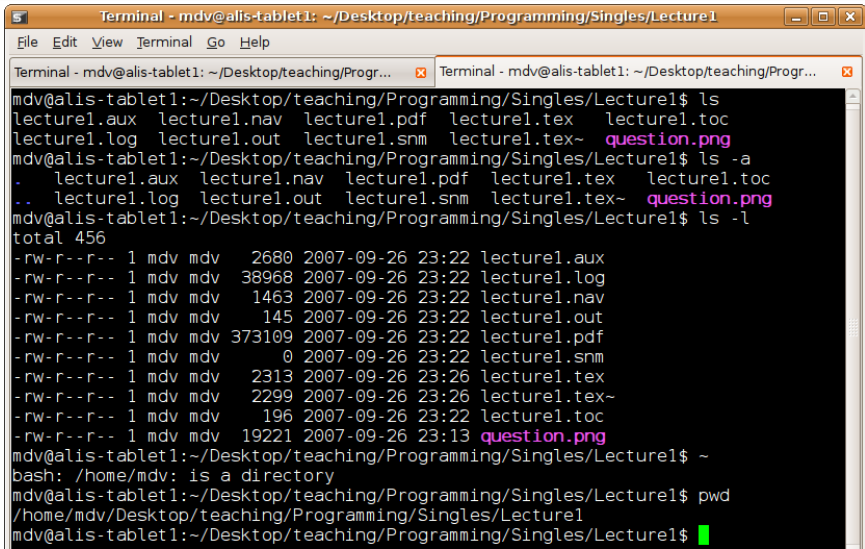
# The Shell

The most important command-line tool is the command shell (often just called the shell)

- Manages a user's interactions with the operating system by:
  - Reading commands from the keyboard
  - Figuring out what programs the user wants to run
  - Running those programs
  - Displaying their output on the screen
- Looks (and works) like an interactive terminal circa 1980



# The Terminal



A screenshot of a Linux terminal window titled "Terminal - mdv@alis-tablet1: ~/Desktop/teaching/Programming/Singles/Lecture1". The window has a menu bar with "File", "Edit", "View", "Terminal", "Go", and "Help". There are two tabs open, both showing the same path. The terminal content shows the following commands and output:

```
mdv@alis-tablet1:~/Desktop/teaching/Programming/Singles/Lecture1$ ls
lecture1.aux lecture1.nav lecture1.pdf lecture1.tex lecture1.toc
lecture1.log lecture1.out lecture1.snm lecture1.tex~ question.png
mdv@alis-tablet1:~/Desktop/teaching/Programming/Singles/Lecture1$ ls -a
. lecture1.aux lecture1.nav lecture1.pdf lecture1.tex lecture1.toc
.. lecture1.log lecture1.out lecture1.snm lecture1.tex~ question.png
mdv@alis-tablet1:~/Desktop/teaching/Programming/Singles/Lecture1$ ls -l
total 456
-rw-r--r-- 1 mdv mdv 2680 2007-09-26 23:22 lecture1.aux
-rw-r--r-- 1 mdv mdv 38968 2007-09-26 23:22 lecture1.log
-rw-r--r-- 1 mdv mdv 1463 2007-09-26 23:22 lecture1.nav
-rw-r--r-- 1 mdv mdv 145 2007-09-26 23:22 lecture1.out
-rw-r--r-- 1 mdv mdv 373109 2007-09-26 23:22 lecture1.pdf
-rw-r--r-- 1 mdv mdv 0 2007-09-26 23:22 lecture1.snm
-rw-r--r-- 1 mdv mdv 2313 2007-09-26 23:26 lecture1.tex
-rw-r--r-- 1 mdv mdv 2299 2007-09-26 23:26 lecture1.tex~
-rw-r--r-- 1 mdv mdv 196 2007-09-26 23:22 lecture1.toc
-rw-r--r-- 1 mdv mdv 19221 2007-09-26 23:13 question.png
mdv@alis-tablet1:~/Desktop/teaching/Programming/Singles/Lecture1$ ~
bash: /home/mdv: is a directory
mdv@alis-tablet1:~/Desktop/teaching/Programming/Singles/Lecture1$ pwd
/home/mdv/Desktop/teaching/Programming/Singles/Lecture1
mdv@alis-tablet1:~/Desktop/teaching/Programming/Singles/Lecture1$
```

# The Shell vs. the Operation System

- The shell is just one program among many
  - Many different ones have been written
  - **sh** was the first for Unix
    - Most others extend its capabilities in various ways
    - Which means that it's the lowest common denominator you can always rely on
  - We will use **bash** (the Bourne again shell)
    - Available just about everywhere
    - Even on Windows (thanks to Cygwin)



# The Shell vs. the Operation System

- In contrast, the operating system is not just another program
  - Automatically loaded when the computer boots up
  - The only program that can talk directly to the computer's hardware
    - I.e., read characters from the keyboard, or send drawing commands to the screen
  - Manages files and directories on the disk
  - Keeps track of who you are, and what you're allowed to do
  - You can run many instances of the shell on a computer at once, but it can only run one operating system at a time





# The File System

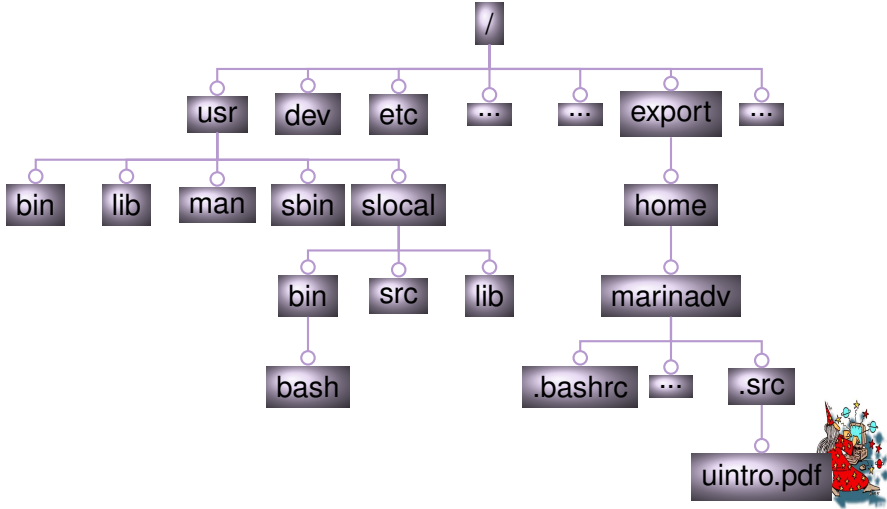
- The file system is the set of files and directories the computer can access

*Everything that stays put when you turn the computer off and restart it*

- Data is stored in files
  - By convention, files have two part names, like notes.txt or home.html
  - Most operating systems allow you to associate a filename extension with an application
    - E.g., .txt is associated with an editor, and .html with a web browser
  - But this is all just convention: you can call files (almost) anything you want
- Files are stored in directories (often called folders)
  - Directories can contain other directories, too
  - Results in the familiar directory tree



# Directory Tree



# Drives

- On Unix, the file system has a unique **root directory** called  
/
  - Every other directory is a child of it, or a child of a child, etc.
- On Windows, every **drive** has its own root directory
  - So C:\home\mdv\notes.txt is different from  
J:\home\mdv\notes.txt
  - When you're using Cygwin, you can also write C:\home\mdv  
as c:/home/mdv
  - Or as /cygdrive/c/home/mdv
    - Some Unix programs give ":" a special meaning, so Cygwin  
needed a way to write paths without it



# Paths

A **path** is a description of how to find something in a file system

- An **absolute path** describes a location from the root directory down
  - Equivalent to a street address
  - Always starts with "/"
  - E.g., /home/mdv is my home directory, and /courses/swc/lec/shell.swc is this file
- A **relative path** describes how to find something from some other location
  - Equivalent to saying, Four blocks north, and seven east
  - E.g., from /courses/swc, the relative path to this file is lec/shell.swc



# Special Paths

- Every program (including the shell) has a current working directory
  - Where am I?
  - Relative paths are deciphered relative to this location
  - It can change while a program is running
- Finally, two special names:
  - "." means the current directory
  - ".." means the directory immediately above this one
    - Also called the parent directory
    - In /courses/swc/data, .. is /courses/swc
    - In /courses/swc/data/elements, .. is /courses/swc/data



# File Systems

Most unix systems have several types of file systems

- Disk-based: UFS : to store all the files users create
- Network-based: NFS: to connect to (mount) drives outside the machine
- tmpfs file system: supports simulating a file system in main memory, possibly backed up by swap storage. This is ideal for temporary files for which fast access is important.
- swap: file system is used to provide backup storage for processes that must temporarily be swapped out
- proc file space: provides a file view on the attributes of processes



# pwd and ls

- **pwd** shows you the current directory

```
pwd
```

```
/home/mdv/Desktop/teaching/Programming/Singles/Lecture1
```

- **ls** shows you what's in the current directory

```
ls
```

```
lecture1.aux  lecture1.out  lecture1.tex  lecture1.vrb  
lecture1.log  lecture1.pdf  lecture1.tex~  question.png  
lecture1.nav  lecture1.snm  lecture1.toc  terminal.png
```



## More on ls

What actually happens when I type ls is:

- The operating system reads characters from the keyboard
- Passes them to the shell (because it's the currently active window on my desktop)
- The shell breaks the line of text it receives into words
- Looks for a program with the same name as the first word (i.e., the command to run)
  - Describe in a moment how the shell knows where to look
- Runs that program
- Reads the program's output and sends it back to the operating system for display





# Flags

- Flags are command-line option you can pass to commands
- Can tell ls to produce more informative output by giving it some flags
- By convention, flags start with "-", as in "-c" or "-l"
- For example: show directories with trailing slash

```
ls -F
```

bluej.png	code.sty~	copyright.tex~	rights.png	uintro.pdf
cm10192.tex	computer.jpg	Doubles/	Singles/	
cm10192.tex~	computer.png	projects.zip	Stylefiles/	
code.sty	copyright.tex*	python.png	template.tex~	

- -a: gives you all files starting with ".", which are normally hidden
- -l: provides long listing format. provides permissions, size, latest access



## Finding your way

- man pages: provide an overview of the functionality of a command.
  - man ls
- apropos: provides all commands related to a certain topic
  - apropos(permissions)
- --help: provides support for a specific command
  - ls --help



# Manipulating Files and Directories

Lets work by example:

- let us create a temporary directory and play around in there

```
mkdir temp
```

- Note: no output
- The -v (verbose) flag tells mkdir to print a confirmation message
- Now go into that directory

```
cd temp
```

- Changes the shell's notion of our current working directory

```
pwd
```

```
/home/mdv/programming1/temp
```



## Manipulating Files and Directories II

- No files there yet:

```
ls -a
```

```
. ..
```

- Use the editor of your choice (emacs,vim) to create a file called earth.txt with the following contents:

```
Name:  Earth
Period: 365.26 days
Inclination: 0.00
Eccentricity: 0.02
Object: Planet
```



## Manipulating Files and Directories III

- Easiest way to create a similar file venus.txt is to copy the one we have

```
cp earth.txt venus.txt
```

```
ls -t
```

```
venus.txt    earth.txt
```

- Note: the -t option tells ls to list newest first
- Check the contents of the file using cat (short for concatenate)
- Just prints the contents of a file to the screen
- You can also use more or less



# Manipulating Files and Directories IV

- Edit the file so that looks like:
- Compare the sizes of the two files using `wc` (for word count) and Compare the two files using `diff`

```
Name: Venus
Period: 224.70 days
Inclination: 3.39
Eccentricity: 0.01
Object: Planet
```

```
wc earth.txt venus.txt
```

```
4   9   69 earth.txt
4   9   69 venus.txt
8  18  138 total
```

```
diff earth.txt venus.txt
```

```
1,4c1,4
< Name: Earth
< Period: 365.26 days
< Inclination: 0.00
< Eccentricity: 0.02
---
> Name: Venus
> Period: 224.70 days
> Inclination: 3.39
> Eccentricity: 0.01
```



# Manipulating Files and Directories V

- Linux does not care about filename extensions.
- `cp earth.txt earth.pdf` is valid although not a very sensible thing to do
- we can rename it using `mv earth.pdf earth2.txt`
- Removing a file can be done using `rm`, like for example `rm earth2.txt`
- A empty directory can be removed with `rmdir` or `rm -r` which recursively removed all files.



# Wildcards

- Some characters (**wildcards**) mean special things to the shell
  - \* matches zero or more characters
    - So `ls *.f77` lists all the Fortran-77 files in a directory

```
wc *.txt
 4   9   69 earth.txt
 4   9   69 venus.txt
 8  18  138 total
```

- ? matches any single character
  - So `ls ???.txt` lists all the text files with two-letter prefixes
  - And `ls ??.*` lists all the files with two-letter prefixes, and any extension
  - on its own means the users home directory
  - `harry` means Harry's home directory
- Note: the shell expands wildcards before running commands
- Note: Be careful using `rm` in conjunction with \*





# Users

- Users have a **user name** and a **password**
- a user also has a **home directory**, and a **shell program**.
- Internally, the system uses so-called **UID** numbers to identify users.
- All this information is stored in the file **/etc/passwd**
- This also stores the user primary group id (GID) identifying a group to which the user belongs.
- A **group** is an arbitrary set of users
- A user can belong to several groups
- `whoami`, `users`, `groups` provide you with information regarding users and groups
- There is one special user with UID 0, called **root**
- This user is often called the super user because he can access all resources on the system, independently of any specific permissions



# Ownership

- Each file has a user as **owner** and a group as **group owner**.
- Using `chmod` the owner can change permissions that determine the type of access (read, write or execute) ...
- allowed to three categories of users: the owner herself, the users belonging to the group owner group, and all other users.
- Note that "execute" permission on a directory is interpreted as "permission to traverse"

```
ls -l
-rw-r--r-- 1 mdv mdv      84 2007-09-27 23:08 earth.pdf
chmod g+w earth.txt
ls -l
-rw-rw-r-- 1 mdv mdv      84 2007-09-27 22:38 earth.txt
```

